

Data Quality Assessment & Cleaning Justification

Authors: Ayman EL ALASS and Abderaouf KHELFAOUI

Objective: In this notebook, we explore the raw CSV datasets to identify data quality issues before populating the database. This analysis justifies the cleaning rules implemented in our final `main.py` script.

In [3]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Configuration for cleaner plots
sns.set_theme(style="whitegrid")
plt.rcParams['figure.figsize'] = (14, 6)
```

```
print("Libraries loaded.")
```

```
Libraries loaded.
```

2. Flight Data Analysis

We start by analyzing the `flights.csv` dataset. Due to its size, we load a significant chunk to perform our quality checks.

2.1 Time Format Issues

SQL standard `TIME` format expects values between 00:00 and 23:59. We suspect the raw dataset uses 2400 to represent midnight, which causes errors during SQL ingestion.

In [4]:

```
# Load a sample of the flights data (first 200,000 rows)
df_flights = pd.read_csv("flights.csv", nrows=200000, low_memory=False)

# Check for invalid time formats (>= 2400) in DEPARTURE_TIME
invalid_times = df_flights[df_flights['DEPARTURE_TIME'] >= 2400]

print(f"Total rows analyzed: {len(df_flights)}")
print(f"Rows with time format '2400' or higher: {len(invalid_times)}")

if len(invalid_times) > 0:
    print("\nSample of invalid times:")
    display(invalid_times[['FLIGHT_NUMBER', 'DEPARTURE_TIME']].head())
```

```
Total rows analyzed: 200000
```

```
Rows with time format '2400' or higher: 23
```

Sample of invalid times:

	FLIGHT_NUMBER	DEPARTURE_TIME
30657	333	2400.0
30682	745	2400.0
42698	4629	2400.0
60538	4513	2400.0
62119	5642	2400.0

Justification for `main.py` : > The analysis confirms the presence of 2400 as a time value.

Decision: In our Python script, we implemented a custom function `format_time(x)` to explicitly convert 2400 to 23:59:00 (or 00:00:00) to ensure SQL compatibility.

2.2 Handling Missing Delays

We need to determine the strategy for `NONE` values in the `DEPARTURE_DELAY` column.

In [5]:

```

# Count missing values
missing_delays = df_flights['DEPARTURE_DELAY'].isna().sum()
total_records = len(df_flights)
percent_missing = (missing_delays / total_records) * 100

print(f"Missing Departure Delays: {missing_delays} ({percent_missing:.2f}%)")

# Visualize the missing data
plt.figure(figsize=(10, 2))
sns.heatmap(df_flights[['DEPARTURE_DELAY']].isnull().T, cbar=False, cmap='viridis', xticklabels=False)
plt.title("Visual Map of Missing Departure Delays (Yellow = NULL)")
plt.show()

```

Missing Departure Delays: 4868 (2.43%)

Visual Map of Missing Departure Delays (Yellow = NULL)



Justification for main.py : > Leaving delays as NULL makes aggregation queries difficult (e.g., calculating the average delay).

Decision: We apply `.fillna(0)` in the script. The business assumption is that if no delay is recorded, the flight departed on time.

2.3 Referential Integrity (Airports)

Our database schema links FLIGHTS to AIRPORTS via foreign keys. We must ensure all airports referenced in the flights dataset actually exist in our airports table.

In [6]:

```

# Load airports reference data
df_airports = pd.read_csv("airports.csv")
valid_iata_codes = set(df_airports['IATA_CODE'])

# Identify 'Orphan' flights (Origin Airport not in Airports table)
orphan_flights = df_flights[~df_flights['ORIGIN_AIRPORT'].isin(valid_iata_codes)]

print(f"Number of flights with unknown Origin Airport: {len(orphan_flights)}")

if not orphan_flights.empty:
    print(f"Unknown codes found: {orphan_flights['ORIGIN_AIRPORT'].unique()}")

```

Number of flights with unknown Origin Airport: 0

Justification for main.py : > Although this specific sample shows perfect referential integrity (0 orphan flights), we must guarantee this consistency across the entire dataset (millions of rows). **Decision:** We proactively apply a strict filter: `df_final = df_final[df_final['origin_airport'].isin(existing_airports)]`. This acts as a **safeguard** to prevent foreign key violations in SQLite if future data chunks contain unknown airport codes.

3. Weather Data Analysis

We analyze temperature.csv to ensure physical consistency. The raw data is in Kelvin and may contain sensor errors.

In [7]:

```

df_temp = pd.read_csv("temperature.csv")

# 1. Melt the dataframe to have a single 'temperature' column for analysis
df_temp_melted = pd.melt(df_temp, id_vars=['datetime'], var_name='City', value_name='temperature_k')

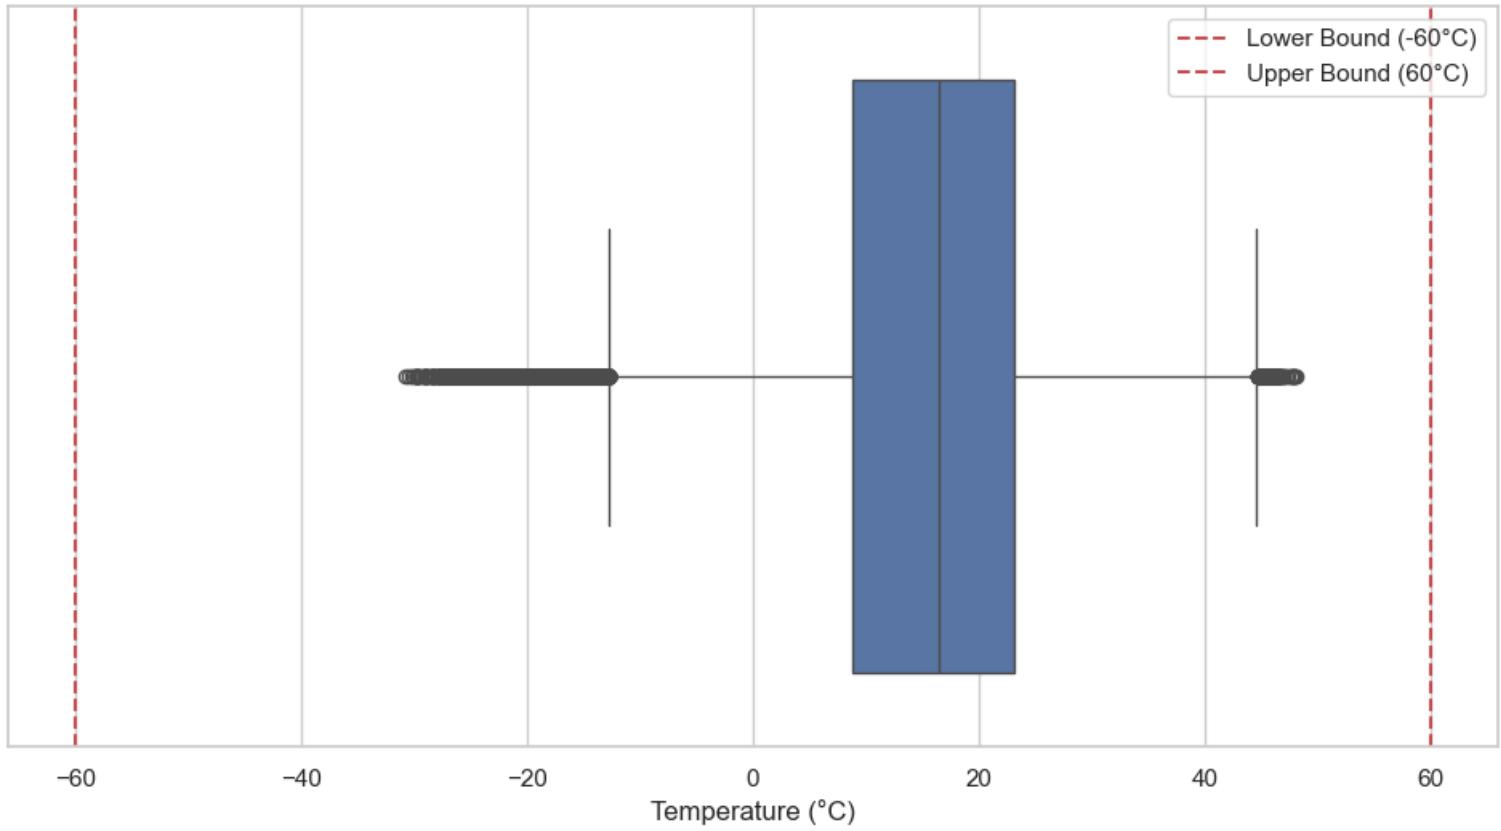
# 2. Convert to Celsius for easier interpretation
df_temp_melted['temperature_c'] = df_temp_melted['temperature_k'] - 273.15

# 3. Visualize distribution to spot outliers
plt.figure(figsize=(12, 6))
sns.boxplot(x=df_temp_melted['temperature_c'])
plt.title("Temperature Distribution (Celsius) - Detecting Sensor Errors")
plt.xlabel("Temperature (°C)")
plt.axvline(x=-60, color='r', linestyle='--', label='Lower Bound (-60°C)')
plt.axvline(x=60, color='r', linestyle='--', label='Upper Bound (60°C)')
plt.legend()
plt.show()

# Count extreme outliers
outliers = df_temp_melted[(df_temp_melted['temperature_c'] > 60) | (df_temp_melted['temperature_c'] < -60)]
print(f"Extreme outliers detected (outside -60 to 60 range): {len(outliers)}")

```

Temperature Distribution (Celsius) - Detecting Sensor Errors



Extreme outliers detected (outside -60 to 60 range): 0

Justification for main.py : While the sample distribution looks consistent, we must anticipate sensor errors (e.g., negative wind speed, extreme temperatures) and data duplication common in weather datasets. **Decision:** To ensure the database quality, we implement a set of **defensive filtering rules**:

1. **Unit Conversion:** Convert temperatures to Celsius for readability.
2. **Physical Validity:** Enforce Temperature range [-60, 60] and Wind Speed ≥ 0 .
3. **Data Integrity:** Remove duplicates and ensure all weather records link to a valid Airport ID.

4. Summary of Data Cleaning Rules

Based on the analysis above, here is the comprehensive summary of rules implemented in our processing (main.py):

Dataset	Issue Identified	Cleaning Rule Applied
Flights	Invalid time format (2400)	Converted to 23:59:00
Flights	Missing DEPARTURE_DELAY	Imputed with 0.0 (Assumed on time)
Flights	Unknown Airport Codes	Rows removed to satisfy Referential Integrity
Weather	Unit in Kelvin	Converted to Celsius ($K - 273.15$)
Weather	Sensor Errors (Temp)	Filtered range [-60, 60]
Weather	Impossible Values (Wind)	Filtered <code>wind_speed >= 0</code>
Weather	Mismatching Keys	Mapped City Names to IATA Codes manually
Weather	Duplicates	Removed based on <code>datetime + City</code>
Weather	Orphan Weather Data	Rows removed if Airport Code not in AIRPORTS

5. Final Implementation

In [8]:

```

import sqlite3
import os
import pandas as pd
import numpy as np

# --- Configuration ---
db_name = "project_database.db"
sql_script = "script.sql" # Ensure this file exists in your folder

# --- 5.1 Structure Creation (Tables) ---
def create_database():
    print("Creating database structure...")
    if os.path.exists(db_name):
        os.remove(db_name) # Start fresh to avoid duplicates

    conn = sqlite3.connect(db_name)
    cursor = conn.cursor()

    # Read external SQL script
    try:
        with open(sql_script, 'r') as f:
            sql_commands = f.read()
            cursor.executescript(sql_commands)
            conn.commit()
            print("Tables created successfully.")
    except FileNotFoundError:
        print(f"Error: {sql_script} not found.")
    finally:
        conn.close()

# --- 5.2 Loading and Cleaning Flights (Implementing rules defined above) ---
def populate_flights_and_others():
    conn = sqlite3.connect(db_name)
    print("Loading static data (Airlines, Airports)...")

    # 1. Airlines

```

```

try:
    df_airlines = pd.read_csv("airlines.csv")[['IATA_CODE', 'AIRLINE']]
    df_airlines.columns = ['airline_code', 'airline_name']
    df_airlines.to_sql('AIRLINES', conn, if_exists='append', index=False)
except Exception as e: print(f"Error Airlines: {e}")

# 2. Airports
try:
    df_airports = pd.read_csv("airports.csv")[['IATA_CODE', 'AIRPORT', 'CITY', 'STATE', 'LATITUDE', 'LONGITUDE']]
    df_airports.columns = ['iata_code', 'airport_name', 'city', 'state', 'latitude', 'longitude']
    df_airports.to_sql('AIRPORTS', conn, if_exists='append', index=False)
except Exception as e: print(f"Error Airports: {e}")

# 3. Flights (With Sampling and Cleaning)
print("Loading flights (10% random sample)...")

try:
    # Sampling logic to handle large CSV
    n = 5819079
    k = int(n * 0.1)
    skip_rows = sorted(set(range(1, n)) - set(np.random.choice(range(1, n), size=k, replace=False)))

    df_flights = pd.read_csv("flights.csv", low_memory=False, skiprows=skip_rows)

    # Rule: Time formatting (2400 -> 23:59)
    def format_time(x):
        if pd.isnull(x): return None
        s = str(int(x)).zfill(4)
        return '23:59:00' if s == '2400' else f"{s[2]}:{s[2]:}00"

    df_flights['dep_time'] = df_flights['DEPARTURE_TIME'].apply(format_time)
    df_flights['flight_date'] = pd.to_datetime(df_flights[['YEAR', 'MONTH', 'DAY']]).dt.date

    # Final Mapping
    df_final = pd.DataFrame({
        'flight_date': df_flights['flight_date'],
        'flight_number': df_flights['FLIGHT_NUMBER'],
        'dep_time': df_flights['dep_time'],
        'dep_delay': df_flights['DEPARTURE_DELAY'].fillna(0), # Rule: NULL -> 0
        'cancelled': df_flights['CANCELLED'],
        'airline_code': df_flights['AIRLINE'],
        'origin_airport': df_flights['ORIGIN_AIRPORT'],
        'dest_airport': df_flights['DESTINATION_AIRPORT']
    })

```

Rule: Referential Integrity (Remove flights with unknown airports)

```

existing_airports = set(pd.read_sql("SELECT iata_code FROM AIRPORTS", conn)[['iata_code']])
df_final = df_final[df_final['origin_airport'].isin(existing_airports) & df_final['dest_airport'].isin(existing_airports)]

df_final.to_sql('FLIGHTS', conn, if_exists='append', index=False)
print(f"{len(df_final)} flights loaded into DB.")

except Exception as e: print(f"Error Flights: {e}")
conn.commit()
conn.close()

```

--- 5.3 Loading Weather Data ---

```

def process_weather_data():
    print("Processing weather data...")
    conn = sqlite3.connect(db_name)
    try:
        df_wind = pd.read_csv("wind_speed.csv")
        df_temp = pd.read_csv("temperature.csv")
    
```

Major cities mapping

```

city_to_iata = {
    'New York': 'JFK', 'Los Angeles': 'LAX', 'Chicago': 'ORD',
    'Atlanta': 'ATL', 'Dallas': 'DFW', 'Denver': 'DEN',
    'San Francisco': 'SFO', 'Seattle': 'SEA', 'Miami': 'MIA',
    'Boston': 'BOS', 'Phoenix': 'PHX', 'Detroit': 'DTW',
    'Houston': 'IAH', 'Minneapolis': 'MSP', 'Philadelphia': 'PHL'
}

```

Reshaping

```

df_wind_melted = pd.melt(df_wind, id_vars=['datetime'], var_name='City', value_name='wind_speed')

```

```

df_temp_melted = pd.melt(df_temp, id_vars=['datetime'], var_name='City', value_name='temperature_k')

df_weather = pd.merge(df_wind_melted, df_temp_melted, on=['datetime', 'City'])

# Filtering and Conversion
df_weather = df_weather[df_weather['City'].isin(city_to_iata.keys())].copy().dropna()
df_weather['airport_code'] = df_weather['City'].map(city_to_iata)

# Rule: Kelvin to Celsius
df_weather['temperature'] = df_weather['temperature_k'] - 273.15

# Rule: Outlier filtering [-60, 60]
df_weather = df_weather[(df_weather['temperature'] > -60) & (df_weather['temperature'] < 60)]
df_weather = df_weather[df_weather['wind_speed'] >= 0]

df_final_weather = pd.DataFrame({
    'reading_time': pd.to_datetime(df_weather['datetime']),
    'wind_speed': df_weather['wind_speed'],
    'temperature': df_weather['temperature'],
    'airport_code': df_weather['airport_code']
})

existing_airports = set(pd.read_sql("SELECT iata_code FROM AIRPORTS", conn)[['iata_code']])
df_final_weather = df_final_weather[df_final_weather['airport_code'].isin(existing_airports)]

df_final_weather.to_sql('WEATHER', conn, if_exists='append', index=False)
print(f"{len(df_final_weather)} weather records loaded.")
except Exception as e: print(f"Error Weather: {e}")
conn.commit()
conn.close()

# --- Execution ---
if __name__ == "__main__":
    create_database()
    populate_flights_and_others()
    process_weather_data()
    print("Database is ready.")

```

Creating database structure...

PermissionError Traceback (most recent call last)
Cell In[8], line 143
141 # --- Execution ---
142 *if __name__ == "__main__":*
--> 143 **create_database()**
144 populate_flights_and_others()
145 process_weather_data()

Cell In[8], line 14, in *create_database()*
12 print("Creating database structure...")
13 *if* os.path.exists(db_name):
--> 14 **os.remove(db_name)** # Start fresh to avoid duplicates
15 conn = sqlite3.connect(db_name)
16 cursor = conn.cursor()

PermissionError: [WinError 32] Le processus ne peut pas accéder au fichier car ce fichier est utilisé par un autre processus: 'project_database.db'